www.editada.org

International Journal of Combinatorial Optimization Problems and Informatics, 18(1), 2026, 17-25. ISSN: 2007-1558.
Doi

Territorial Information Retrieval from Heterogeneous Open Data: Extending a Water Data Warehouse from Mexico City to its Metropolitan Area

Eduardo Uriel Velázquez Arrieta[1], Omar Fernando Pulido Morales[1], Emilio García López[1], Clara Aide Hernández Martínez[1] and Gabriel Hurtado Avilés[1]

[1] Escuela Superior de Cómputo (ESCOM), Instituto Politécnico Nacional, Mexico City, Mexico

{evelazqueza1900, opulidom2000, egarcial2001, chernandezm1903}@alumno.ipn.mx
ghurtadoa@ipn.mx
velazquez.arrieta.uriel@gmail.com, pulido.morales.omar@gmail.com
egarlo@hotmail.com, claraahm11@gmail.com, gabrielhuav@gmail.com

✉Corresponding author: gabrielhuav@gmail.com

Abstract. This work presents a data warehouse for the territorial information retrieval of water consumption in Mexico City, built from heterogeneous open sources and extended to its metropolitan area. The contribution lies in three elements: an ETL process that consolidates SACMEX data into a dimensional model of 70,886 records, a declarative R2RML mapping that republishes the star schema as an RDF knowledge graph, and a deployment methodology that guarantees a demonstration through a static version. Against the 2020 Census, billed volumes amount to 47 litres per inhabitant per day, 38 % of published consumption, least in the periphery, and the atypical-consumption score structurally misses drops. The extension integrates 62 verified files on 63 municipalities into 2.76 million rows: municipal volumes require a quality filter, and neighbourhoods with almost no billed volume report shortages nearly as often as the rest. Every layer of a new map is a verified SQL query.

Keywords: Data warehouse, Open data, Knowledge graph, RDF Data Cube, Water consumption, Anomaly detection

Article information
Received:
Accepted:

1 Introduction

Mexico City faces a water crisis of aquifer overexploitation, network leaks and population growth. Open data on water consumption exist (SACMEX), but in poorly accessible formats and without query tools, which limits their use for evidence-based decision making.

This work addresses the problem from two perspectives proper to information processing: the extraction of structured information from heterogeneous open sources through an ETL process, and the retrieval of that information through a parameterized query interface, transforming dispersed datasets into a queryable dimensional repository that anyone can filter and visualize geospatially by neighbourhood and borough.

The contribution of this work is threefold: (i) a reproducible ETL process to consolidate official open data into a dimensional data warehouse, (ii) a declarative R2RML mapping that republishes it as a validated RDF knowledge graph, and (iii) a continuous-deployment methodology that keeps a functional demonstration always accessible, combining a dynamic deployment connected to a database with a static version published in a public repository.

The analysis of the 70,886 consolidated records reveals that the boroughs of Cuauhtémoc and Miguel Hidalgo lead the total consumption (Table 3), and that consumption varies across the available bimesters without a direct causal relationship with dry or rainy periods. The system serves citizens querying their own neighbourhood, public servants diagnosing territorial demand, researchers, journalists and civil organisations monitoring water management.

This article extends the version presented at ICOKG 2026 (Velázquez Arrieta et al., in press). It adds five elements. The first is an account of what a consumption value in the source actually is, taken from the publisher's clarifying note (Section 3.4). The second measures, against the 2020 Census, how much of the city's consumption the billed volumes in the export represent and how the borough ranking changes per inhabitant (Section 5.3). The third is a sensitivity analysis of the atypical-

© Editorial Académica Dragón Azteca S. de R.L. de C.V. (EDITADA.ORG). All rights reserved.

consumption module over the prevalence, composition and magnitude of the anomalies, which shows that the deployed statistic is structurally blind to consumption drops (Section 5.5). The fourth is a diagnosis of the same statistic on the real data, with a log-scale variant and the evidence that each scale is biased in one direction (Section 5.6). The fifth extends the warehouse from the city to its metropolitan area and publishes the result in a map whose layers are verified SQL queries (Section 5.7).

Section 2 reviews related work, Section 3 describes the methodology, Section 5 presents the results and Section 7 the conclusions.

2 Related Work

2.1 Foundations: Data Warehouses, Dimensional Modelling, and ETL

Dimensional modelling and the star schema constitute the consolidated basis for the design of analytical data warehouses. Kimball and Ross (2013) establish the canonical principles —fact and dimension tables, surrogate keys, slowly changing dimensions, and the ETL subsystems— that ground the model adopted in this work. On the integration phase, Vassiliadis (2009) offers a reference survey of *Extract–Transform–Load* technology, covering the conceptual and logical modelling of the processes and the cross-cutting extraction, transformation, and loading issues that methodologically support our pipeline.

2.2 Data Warehouses and Information Systems for Water Management

In the water domain, the most direct antecedent is that of Soares et al. (2019), who implement a data warehouse with an ETL process to monitor water supply and consumption, integrating meter readings, reservoir volumes, and geographic location from management-application databases and heterogeneous files. Abecker et al. (2014) present a hybrid sensor-semantic architecture (the WatERP project) that combines a relational/PostGIS store with a triple store for integrated resource management. These systems, however, mostly start from operational data, SCADA, or proprietary sensor networks, and not from open government data.

2.3 Water Consumption in Mexico City from Utility Records

Peer-reviewed spatial analyses of consumption in Mexico City come mainly from two groups, and both obtained their data from the utility rather than from the open portal. Medina-Rivas et al. (2022) analyse the 2019 records, the year of the open export used here, report a domestic consumption of 123 litres per person per day, and find two clusters of unequal consumption explained by discontinuous supply, dependence on external sources, altitude and socioeconomic status; a later study finds that per-capita domestic use fell by 40 % between 2010 and 2019 while its spatial distribution became more unequal (Medina-Rivas et al., 2025). Ramos-Bueno et al. (2021) obtained colonia-level consumption and metering for 2008–2013 through public-information requests and explain residential demand with spatial regression, after excluding colonias with zero consumption or no metering coverage, and a panel study by the same group extends the series to 2004–2022 (Ramos-Bueno et al., 2024). These works answer substantive questions with data a reader cannot obtain from the portal. The present one addresses the complementary problem of making the open export itself queryable and verifiable, and what it finds about the billing basis and about colonias without billed volume (Sections 3.4 and 5.6) is consistent with the exclusions those studies apply.

2.4 The Gap: ETL over Heterogeneous Open Data and Data Quality

Integrating open government data into dimensional warehouses poses its own challenges. Berro, Megdiche, and Teste (2015) propose “content-driven” ETL processes that address precisely the semantic and structural heterogeneity, the absence of schemas, and the dispersion characteristic of open data, factors that overwhelm traditional ETL. Complementarily, Vetrò et al. (2016) define a quantitative measurement framework for the quality of Open Government Data (accuracy, completeness, consistency, traceability, and freshness) and empirically report that the most frequent problems are missing metadata, incomplete data, and the absence of update information —exactly the challenges our ETL faces over the SACMEX data.

2.5 Geospatial Information Retrieval and Visualization

Purves et al. (2018) offer a reference monograph on Geographic Information Retrieval (GIR) —indexing, search, and spatial ranking— that theoretically frames the territorial information-retrieval component of this work.

2.6 Knowledge Graphs for Urban and Water Domains

Republishing a relational warehouse as RDF is a mature practice with W3C standards for both ends. The RDF Data Cube vocabulary (World Wide Web Consortium, 2014) models statistical observations over conformed dimensions, which is the same structure a star schema encodes; R2RML (World Wide Web Consortium, 2012) specifies the mapping declaratively, so the graph is derived rather than maintained by hand; and GeoSPARQL 1.1, documented by Car and Homburg (2022), supplies the geometry model and the topological vocabulary that a flat territorial dimension cannot express. In the water and urban domains the direction is established: Roussey et al. (2023) survey knowledge-graph technologies as the next frontier for food,

agriculture and water. The closest graph in the water domain is the Water Health Open Knowledge Graph (Lippolis et al., 2025), built in a European project, which models water consumption together with pollution, extreme weather, infectious disease and drug distribution; we found no RDF knowledge graph built over Mexican utility water data. What these works share is the premise this article adopts: the graph is valuable for interoperability and topological querying, not as a replacement for the analytical store.

2.7 Reproducibility and Deployment of Scientific Demonstrations

Finally, Sandve et al. (2013) formulate widely cited rules for computational reproducibility (version control, documentation of preprocessing, public access to data and scripts) that ground our reproducible static version. Hill et al. (2023) provide a practical guide to developing and deploying an open-source, reproducible web dashboard for water-sanitation data, a close domain that supports the continuous-deployment methodology with reproducible artefacts.

2.8 Positioning of the Contribution

Unlike existing water data warehouses —which start from operational data, SCADA, or proprietary sensors (Abecker et al., 2014; Soares et al., 2019)— and recognizing that the integration of heterogeneous open data is a documented but unsolved problem in this domain (Berro et al., 2015; Vetrò et al., 2016), this work differs in three respects. First, an ETL process designed to consolidate official SACMEX open data into a dimensional star schema. Second, the exposure of the information through parameterized queries and a choropleth geospatial visualization that enables territorial information retrieval (Purves et al., 2018), published as static artefacts that require no server. Third, a continuous-deployment methodology accompanied by a reproducible static version that guarantees the indefinite verifiability of the system (Hill et al., 2023; Sandve et al., 2013). This same approach of consolidating official open data was previously applied to the seismic-risk domain in Mexico (Villa Vargas et al., 2026), a direct methodological antecedent transferred here to the water domain. To the best of our review, no academic, peer-reviewed, and reproducible data warehouse built specifically over SACMEX/CDMX open data exists.

3 Methodology

3.1 General Architecture

The system follows a layered architecture: a data layer (PostgreSQL), a processing layer (ETL with Python and SQL), a query layer of parameterized SQL and, over it, a declarative RDF projection, and a presentation layer (HTML/CSS/JS frontend with Leaflet and Plotly). The warehouse implements a star schema over two fact tables —`fact_consumo_agua`, the subject of this article, and `fact_clima`, the climate series joined through the shared temporal dimension— and three conformed dimensions: `dim_tiempo` (date, year, bimester), `dim_ubicacion` (borough, neighbourhood, latitude, longitude) and `dim_indice_des` (development level). Identifiers are quoted verbatim from the published schema.

3.2 ETL Consolidation Process

The first contribution of the work is the ETL process that consolidates the SACMEX open data into the dimensional warehouse. The complete detail of the process —the DDL and the ETL scripts (`ddl/0_schema.sql`, `etl/1_copy.sql`, `etl/2_dim.sql`, `etl/3_fact.sql`) and the parameterized retrieval queries— is documented in the project's public repository; this article presents only the architecture of the process.

The pipeline operates in four phases. (1) **Extraction**: reading the open SACMEX CSV files. (2) **Validation and cleaning**: discarding records with empty fields, inconsistent coordinates, or incompatible formats. (3) **Loading to staging and populating dimensions**: `COPY` of the data into temporary tables and extraction of the unique values into `dim_tiempo`, `dim_ubicacion` and `dim_indice_des`. (4) **Populating the fact tables**: a join between staging and dimensions that populates `fact_consumo_agua` and `fact_clima`, finally dropping the staging tables. The information thus consolidated becomes available for retrieval through the parameterized queries distributed with the archived release.[1]

3.3 Continuous Deployment Methodology

So that the demonstration survives without a server, the system is published in two modalities. The **dynamic deployment** keeps the complete architecture over PostgreSQL, serves parameterized queries and produced the results reported here. The **static version** exports the data as `JSON` and `GeoJSON` that the frontend consumes directly, so querying, filtering, aggregation and visualization run on the client. The duality addresses the recurring problem that academic prototypes stop working shortly after evaluation (Hill et al., 2023; Sandve et al., 2013). The system is available at https://github.com/gabrielhuav/Data_Warehouse_static (repository), https://gabrielhuav.github.io/Data_Warehouse_static/ (static demo, no backend required) and https://github.com/gabrielhuav/data_warehouse_cdmx (warehouse, preservation fork).

[1] SACMEX operates a proprietary platform integrating SCADA and multiple sources; it is distinguished from a reproducible academic work over open data such as this one.

What is preserved deserves precision. The archived release (https://github.com/gabrielhuav/Data_Warehouse_static/releases/tag/v1.5-icokg2026) distributes the warehouse itself —schema, ETL, both source files and the container definition that rebuilds the database, vendored unmodified from its author's repository and additionally preserved as a fork— together with the static frontend, the R2RML mapping, a browser-side SPARQL explorer (https://gabrielhuav.github.io/Data_Warehouse_static/grafo.html) and the scripts that regenerate the exported data, so the JSON the demonstration consumes is an export of the warehouse and not a separate dataset. Citations refer to the frozen release; the repository head may evolve, the release will not.

3.4 Data Scope, Grain and Cleaning Criteria

Four properties of the source bound every result reported below. *Extent*: the resource is published as *Consumo Agua (1er Semestre 2019)* (SACMEX, 2026) and contains exactly that —the first three bimesters of 2019, one reading date each, in 71,102 records— so the warehouse is not a multi-year series and the seasonal variation it supports is a within-semester comparison. The bound is declared by the publisher in the title of the resource, not chosen by us, and when the portal migrated the resource in August 2026 it still held those three bimesters, identical in content to the file consolidated here (same MD5, `41495b03`…, once line endings are normalised): a partial extraction published once, without versioning, which is the problem this article studies. Against the source becoming unavailable, both files the ETL consumes are redistributed with the archived release and the portal page is preserved in the Internet Archive. *Grain*: one fact row is one *manzana* (city block) for one billing period, and the export carries 22,930 distinct coordinate-bearing units over 1,553 borough–neighbourhood pairs, so the location dimension is coarser than the grain and every aggregate is a roll-up. *Scale*: the development index has four ordinal levels —ALTO, MEDIO, BAJO and POPULAR— published with the dataset rather than derived here.

Billing basis: a consumption value is a billed volume, not a meter reading. The clarifying note that the publisher distributes with the data dictionary, grounded in the 2019 Fiscal Code of Mexico City, states that an intake is billed by its meter where one is installed and authorised, and otherwise by an average: the average of its cadastral *colonia* when at least 70 % of the intakes there are metered, or the average of the intake's last six metered bimesters when its meter is broken or cannot be read. The export does not say which basis applies to each record, so part of every aggregate is imputed by construction, in a proportion the data do not reveal. Two consequences follow. A neighbourhood with unusually uniform consumption may reflect averaging rather than uniform behaviour, and a detector run on these values ranks departures in billing, which coincide with departures in physical use only where intakes are metered. The same note defines the *manzana* as a cadastral unit, identified by digits of the cadastral account number rather than by the census geography, so joining the warehouse to census data requires a spatial join on the coordinates, not a key match; and it defines the development index as a statistical construct over socioeconomic variables aggregated to the block, which is why the knowledge graph treats the four levels as an ordinal scale rather than as measured quantities.

The exclusion criteria are deterministic and enforced structurally: the fact table is populated by inner joins against the three dimensions, so a record that fails to resolve a dimension member is not loaded. Table 1 quantifies each criterion rather than merely declaring data quality (Vetrò et al., 2016). The distribution is unusually clean and that is itself a finding: a single criterion accounts for every discarded record —the 216 rows carrying the literal `NA` in both territorial fields— and none is lost to a malformed measure or an unparseable date. This is a property of the SACMEX export rather than a merit of the process.

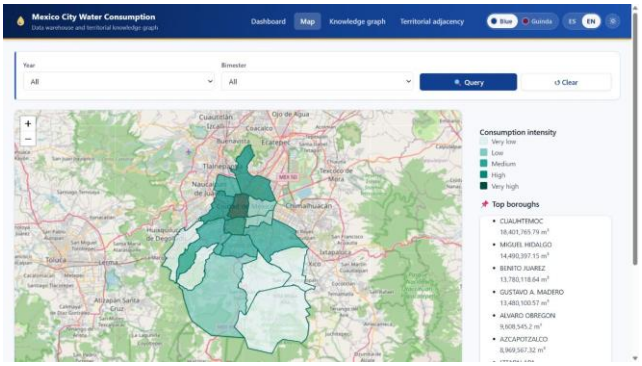
Table 1. Records discarded during validation and cleaning, over the 71,102 records of the source export.

Exclusion criterion	Records	% of source
Unmatched territorial label	216	0.30
Missing development index	0	0.00
Reading date absent from the time dimension	0	0.00
Total discarded	216	0.30
Loaded into the fact table	70,886	99.70

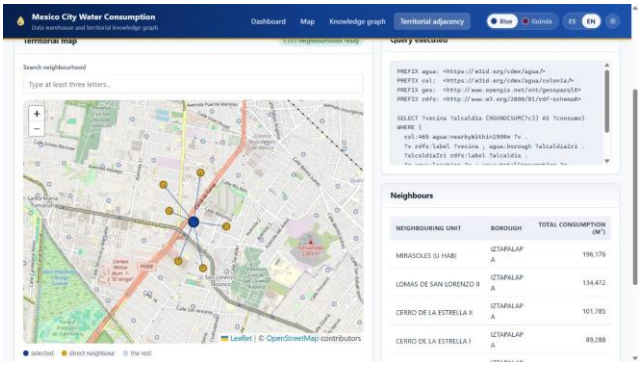
Data sources and attribution. The meteorological series comes from Open-Meteo under CC BY 4.0 and the cartographic base of the interactive map is OpenStreetMap under the Open Database License, which requires attribution wherever the map is reproduced.

3.5 Frontend and Geospatial Visualization

The frontend offers a dashboard with indicators, interactive charts (Plotly), and a choropleth map of the 16 boroughs (Leaflet + GeoJSON), where colour represents the consumption level and each borough displays a popup with its detail (Fig. 1a). Alongside it, and over the same exported data, a second view resolves territorial adjacency by traversing the knowledge graph of Section 4 rather than by filtering the cube: selecting a neighbourhood returns its neighbours ranked by consumption, together with the `agua:nearbyWithin1500m` query that produced them (Fig. 1b). The two views in Fig. 1 are therefore the same warehouse read through its two representations, and both run in the browser with no server behind them.



(a) Choropleth retrieval by borough.
https://gabrielhuav.github.io/Data_Warehouse_static/mapa.html



(b) Adjacency as graph traversal, with the `agua:nearbyWithin1500m` *query and its result.*
https://gabrielhuav.github.io/Data_Warehouse_static/vecindad.html?alcaldia=IZTAPALAPA&colonia=AO%20DE%20JUAREZ

Fig. 1. The two retrieval modes of the static demonstration, both running in the browser without a server; each caption links to the exact view captured. Base map © OpenStreetMap contributors (ODbL).

3.6 Implementation of the Atypical-Consumption Module

The current version replaces supervised prediction with an exploratory detection module oriented to analytical retrieval: because the open source carries no anomaly labels, the system computes statistical indicators over the retrieved records instead of training a model. The procedure runs unsupervised on the frontend. Consumption is aggregated by neighbourhood, each neighbourhood is compared against the others of its own borough, and the units with the largest deviation are ranked as candidates for inspection. Flagging patterns that require later review keeps the practical value of the component without presenting a prediction as if it had been trained against ground truth.

Formally, each neighbourhood g receives an atypicality score:

$$A_g = \left| \frac{x_g - \mu}{\sigma + \varepsilon} \right|, \tag{1}$$

where x_g is the aggregate consumption of a neighbourhood, μ and σ are the mean and standard deviation of the neighbourhoods of its own borough —the comparison set implemented in `assets/atipicos.js`—, and ε prevents division by zero. The groups with the largest A_g are reported in the interface as zones with atypical consumption.

Two clarifications delimit the module. It is *exploratory* and not a causal predictor: a high A_g states that a zone departs from its comparison set, not why. And it does *not* perform real-time leak detection, since no leak reports, work orders or network telemetry are integrated into the warehouse and the evaluation in Section 5 rests exclusively on synthetically injected perturbations. Its role is the ranking function of the retrieval framework: the choropleth answers “how much is consumed here”, the module answers “which units should I look at first”, which is the ordering problem at the centre of information retrieval (Purves et al., 2018).

4 From the Dimensional Model to a Knowledge Graph

4.1 Rationale: Why a Star Schema, and What a Graph Adds

The star schema was not chosen by default. The dominant workload is the repeated aggregation of one numeric measure over three conformed dimensions with filters chosen interactively, which is what dimensional modelling was designed for (Kimball & Ross, 2013): constant join depth and aggregation by a single grouped scan. A graph would pay a traversal cost for joins that foreign keys already resolve. What a graph adds is complementary: *interoperability*, since a standard vocabulary makes the resource linkable to other municipal datasets without prior schema agreement, the recurring obstacle in the open-data

literature (Vetrò et al., 2016); *topological expressiveness*, since questions such as "which neighbourhoods adjacent to an atypical zone also exceed the borough mean" are natural traversals and awkward over a flat territorial dimension; and *entity reconciliation*, since stable identifiers would resolve the label heterogeneity our ETL absorbs through normalisation rules — the `w3id` namespace used here is not yet registered for dereferencing, which we leave as future work. We therefore keep the star schema as the analytical store and derive the RDF layer by a declarative mapping, so that neither representation is maintained by hand.

4.2 Vocabulary Alignment and R2RML Mapping

The correspondence between a star schema and the RDF Data Cube (World Wide Web Consortium, 2014) is close to mechanical, which is what makes this layer inexpensive to maintain: a fact row becomes a `qb:Observation`, the measures become `qb:MeasureProperty`, the time and location dimensions become `qb:DimensionProperty`, the neighbourhood entity becomes a `geo:Feature` with its `geo:asWKT` geometry, territorial adjacency becomes the project-defined `agua:nearbyWithin1500m`, and the development index becomes a `skos:Concept` in a declared `skos:ConceptScheme` with an explicit numeric rank. The mapping is expressed in R2RML (World Wide Web Consortium, 2012) and distributed with the archived release, declaring seven triples maps over both fact tables, the territorial dimension and its geometry, the derived adjacency, the temporal dimension and the index scheme. A companion `schema.ttl` supplies what the mapping alone cannot: the `qb:DataStructureDefinition` for both cubes, its components typed as dimension, measure or attribute, the cubic-metre unit and the ordinal scale. Loaded together with the browser graph, it satisfies the Data Cube integrity constraints IC-1, IC-2, IC-3 and IC-11 with zero violations, reproducible with `scripts/validar_datacube.py`; the check is stated for the published demonstration graph, not for the complete materialisation. The alignment is therefore a validated property rather than a nominal reuse of two terms. Materialising it yields on the order of 7.6×10^5 triples over 71,067 observations, 1,553 territorial features, 181 temporal entities and the four index members.

Two properties of the geospatial layer bound what the traversal can claim. The warehouse stores *points*, not polygons: each neighbourhood is the centroid of its readings, recovered from coordinates present in the source export but dropped by the original loading schema. The substitution is defensible at this scale —the 22,930 distinct coordinates are never shared between two neighbourhoods and the median neighbourhood diameter is 0.63 km— though coarse for the ten neighbourhoods above three kilometres. Because centroids are points, no Simple Features boundary relation applies: `geo:sfTouches` is undefined for point/point geometries, and relabelling a distance threshold as a standardised topological predicate would misstate what the data support. The mapping therefore declares its own `agua:nearbyWithin1500m`, documented in `schema.ttl` as centroid proximity below 1.5 km. Official polygons would allow `geo:sfTouches` in the strict sense; that is a substitution in the source data, and we leave it as future work.

4.3 Territorial Retrieval as Graph Traversal

The value of the layer is that it makes a class of questions expressible that the tabular interface does not naturally support. GeoSPARQL supplies the geometry model (Car & Homburg, 2022); Listing 1 is the query the published interface executes verbatim, retrieving the neighbourhoods near a selected zone together with their own consumption, a single traversal that the relational interface resolves only through an ad-hoc spatial self-join, and which Fig. 1b shows executing in the published demonstration.

Listing 1. Territorial traversal over the knowledge graph, as executed by `vecindad.html`.

```
PREFIX agua: <https://w3id.org/cdmx/agua/>
PREFIX col:  <https://w3id.org/cdmx/agua/colonia/>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>

SELECT ?vecina ?alcaldia (ROUND(SUM(?c)) AS ?consumo) WHERE {
  col:465 agua:nearbyWithin1500m ?v .
  ?v rdfs:label ?vecina ; agua:borough ?alcaldiaIri .
  ?alcaldiaIri rdfs:label ?alcaldia .
  ?o agua:location ?v ; agua:totalConsumption ?c .
} GROUP BY ?vecina ?alcaldia ORDER BY DESC(?consumo)
```

That the graph is queryable, and not merely declarable, is shown without any server: the static deployment ships a browser-sized subset of 140,499 triples with a SPARQL 1.1 engine vendored in the release itself, so the demonstration depends on no external service. The subset is an aggregate, not a slice: its observations live at the neighbourhood-period grain in a separate `agua:aggobs/` namespace under `agua:consumoCDMXDemo`, linked to the complete dataset through PROV-O. Distinct identities avoid the contradictory RDF that reusing an IRI across grains would produce if the graphs were merged; the price is that a query written against the subset is not guaranteed to run unchanged against the full materialisation. The two

representations agree where they overlap: aggregating consumption by bimester through SPARQL returns the same totals the reproduction script obtains in SQL.

5 Results and Evaluation

5.1 Integrated Data Volume

The ETL process consolidated the water-consumption information of Mexico City into the dimensional warehouse. Table 2 presents the stored volume. Each cardinality answers to a different grain and this must be read carefully. The 1,553 locations are the distinct borough–neighbourhood pairs and are coarser than the fact grain. The four index levels are the four members of the SACMEX ordinal scale, not three levels plus a residual. And the 181 time records are *not* reading dates: they are the daily observations of the climate series integrated alongside the consumption data, from which the dimension is populated. The consumption source carries three reading dates only, one per bimester, so the fact table distributes its rows over those three while the climate facts use all 181. The shared temporal dimension at daily grain is what allows both to be queried together.

Table 2. Data volume stored in the data warehouse

Table	Records
fact_consumo_agua	70,886
dim_ubicacion	1,553
dim_tiempo	181
dim_indice_des	4

5.2 Findings of the Consumption Analysis

The analysis of the 70,886 consolidated records reveals the following consumption patterns in Mexico City:

- **Territorial concentration:** The boroughs with the highest total consumption are Cuauhtémoc and Miguel Hidalgo, followed by Benito Juárez and Gustavo A. Madero (Table 3).
- **Temporal distribution:** Consumption shows variations across the bimesters available in the warehouse. In the aggregated query by bimester, the highest total consumption was observed in bimester 2, followed by bimester 3 and bimester 1. Therefore, no direct causal relationship with dry or rainy periods is established without integrating additional climatic variables.
- **Relationship with the development index:** Neighbourhoods with a High development index show a higher average consumption than those with a Low index, a descriptive association discussed in Section 7.

Table 3. Five boroughs with the highest total water consumption

Borough	Total consumption
Cuauhtémoc	18 401 765.79
Miguel Hidalgo	14 490 397.15
Benito Juárez	13 780 118.64
Gustavo A. Madero	13 480 100.57
Álvaro Obregón	9 608 545.20

5.3 Billed Volume and Consumption per Inhabitant

The conference version warned that totals are not a per-capita indicator (Section 6). Dividing the domestic billed volume of the semester by the residents of private dwellings counted in the 2020 Census, the nearest count to 2019 (Instituto Nacional de Estadística y Geografía, 2021), and by the semester's 181 days gives the figures of Table 4. Two results follow. The first

is that the ranking changes: Benito Juárez, third in total consumption, is first per inhabitant with 126.5 litres per day, while Gustavo A. Madero and Iztapalapa, fourth and seventh in totals, fall to tenth and fourteenth. The second matters more. The city-wide figure is 47.3 litres per inhabitant per day, or 53.9 including mixed-use billing, about 38 % of the 123 litres per day that Medina-Rivas et al. (2022) report for the same year from records provided by the utility. The gap does not come from the export being incomplete. Annualised, its billed volume is about 243 million cubic metres, close to the 255 million that the utility reported as billed in 2021 in answer to a public-information request (INFO CDMX, 2022). It is billing that captures only part of the water used: the export amounts to 7.7 m^3/s, about half of the 15.44 m^3/s of total consumption that Ramos-Bueno et al. (2024) estimate for the city. The shortfall is concentrated in the periphery, where 16.0 litres in Iztapalapa, 8.2 in Tláhuac and 2.5 in Milpa Alta cannot be the water people use. The clarifying note covers metered and average billing only (Section 3.4), so intakes billed without a volume would leave no trace in the data; the Fiscal Code applies a fixed fee rather than a volume to domestic intakes in colonias whose supply is officially classified as by rota (*tandeo*). That explanation is plausible but cannot be confirmed from the export alone. Either way, the open export is a faithful image of volumetric billing rather than of consumption, and Table 3 ranks billing coverage as much as demand. This is a question of validity rather than of completeness in the sense of open-data quality (Vetrò et al., 2016): every record may be present and the measure still answer a different question from the one its title suggests. The table is reproducible with `scripts/per_capita_alcaldia.py`.

Table 4. Domestic billed volume of the first semester of 2019 per resident of private dwellings (2020 Census), by borough, with the change in rank from total consumption (all uses) to consumption per inhabitant.

Borough	Residents	Domestic (thousand m^3)	Litres per inhabitant per day	Rank: total to per inhabitant
Benito Juárez	432,259	9,894.9	126.5	3 to 1
Miguel Hidalgo	412,529	8,640.5	115.7	2 to 2
Cuauhtémoc	540,976	10,207.7	104.2	1 to 3
Azcapotzalco	431,347	6,153.8	78.8	6 to 4
Iztacalco	403,805	4,663.5	63.8	10 to 5
Venustiano Carranza	443,120	5,007.0	62.4	9 to 6
Cuajimalpa de Morelos	217,606	2,064.3	52.4	12 to 7
Álvaro Obregón	757,117	6,590.2	48.1	5 to 8
Coyoacán	612,130	5,128.3	46.3	8 to 9
Gustavo A. Madero	1,162,914	9,196.4	43.7	4 to 10
Tlalpan	695,308	2,843.1	22.6	11 to 11
La Magdalena Contreras	246,785	845.0	18.9	14 to 12
Xochimilco	438,038	1,304.1	16.4	13 to 13
Iztapalapa	1,820,809	5,263.0	16.0	7 to 14
Tláhuac	391,967	583.0	8.2	15 to 15
Milpa Alta	152,682	69.1	2.5	16 to 16
Mexico City	**9,159,392**	**78,453.9**	**47.3**	

5.4 Quantitative Evaluation of the Atypical-Consumption Module

With no anomaly labels in the open source, the module is evaluated against injected ground truth. The protocol scores *the deployed function itself*, not a proxy: consumption is aggregated by neighbourhood, μ and σ are computed over the neighbourhoods of each borough, and A_g ranks them, which is exactly what `assets/atipicos.js` computes. The synthetic set reproduces the structure of the warehouse —1,553 borough–neighbourhood pairs over three bimesters— with 47 injected anomalies per run, and the experiment is repeated over 20 seeds. The three detectors are compared at equalised operating points, each emitting as many alerts as there are injected anomalies, so precision, recall and F1 coincide by construction and the areas under the curves characterise the ranking independently of any threshold (Table 5).

Table 5. Detection of injected anomalies, mean ± standard deviation over 20 seeds at equalised operating points. Reproducible with `evaluacion-ml/eval_anomalias.py`, a script reconstructed for this artefact.

Method	P = R = F1	ROC-AUC	PR-AUC
A_g (deployed)	0.581 ± 0.040	0.966 ± 0.013	0.677 ± 0.047
Isolation Forest	0.650 ± 0.053	0.971 ± 0.013	0.728 ± 0.058
LOF	0.493 ± 0.075	0.846 ± 0.036	0.464 ± 0.075

The deployed statistic is competitive with the strongest baseline: its ROC-AUC of 0.966 lies within one standard deviation of the 0.971 of Isolation Forest (Liu et al., 2008), and it recovers 58 % of the injected anomalies against 65 %. That 0.069 difference is modest relative to the combined between-seed variability, although Isolation Forest remains consistently superior at the evaluated operating point. LOF (Breunig et al., 2000) is clearly worse (0.846), since it looks for locally sparse neighbourhoods whereas the injected perturbations are departures in global magnitude. An interpretable borough-relative z-score therefore loses little ranking quality against an ensemble method a municipal analyst could not inspect, which supports A_g as a conservative exploratory filter, one that ranks a short list a utility can afford to inspect, and point to Isolation Forest as the natural improvement.

Two limitations remain. The perturbations are synthetic and their magnitude is chosen, so the figures measure sensitivity to known deviations rather than field performance against real leaks, for which the system is not designed and no labelled data exist. And the comparison rests on a single injection model; a different anomaly shape could reorder the methods.

5.5 Sensitivity of the Detection Module

The second limitation can be removed experimentally. The protocol was repeated varying one factor of the injection model at a time: the prevalence (1, 3, 5 and 10 %), the share of spikes among the anomalies (10 to 90 %) and their magnitude, from subtle (spikes ×1.5–2, drops ×0.50–0.67) through moderate (×2–3 and ×0.33–0.50) to the strong setting of Table 5 (×3–6 and ×0.05–0.25). Each of the twelve settings ran over the same 20 seeds, 240 runs in all, with the same territorial identities, deployed function, baselines and equalised operating point; recall is additionally reported separately for spikes and for drops. With the published setting the script reproduces Table 5 exactly, and it adds a fourth detector, log-A_g: the same statistic computed on the logarithm of the aggregate consumption. Because anomalies are at most a tenth of the units, the area under the precision–recall curve is the more informative summary of a ranking (Saito & Rehmsmeier, 2015). Everything is reproducible with `evaluacion-ml/sensibilidad_anomalias.py`.

The deployed statistic turns out to be structurally blind to drops (Fig. 2a). Under the published setting it finds 82 % of the injected spikes and 2 % of the drops. The cause is arithmetic, not statistical. A drop can lower a neighbourhood's consumption by at most its whole value, so a neighbourhood below its borough mean scores at most μ/σ, the inverse of the borough's coefficient of variation, whereas a spike's score is unbounded. With the dispersion of the synthetic boroughs that ceiling is about 2.3, below the scores that ordinary neighbourhoods in the right tail already reach, so when the alert budget is spent on the largest scores the drops are never reached. The published precision of 0.58 is therefore the share of spikes found weighted by the share of spikes injected, and it falls to 0.12 when nine anomalies in ten are drops (Fig. 2b and Table 6).

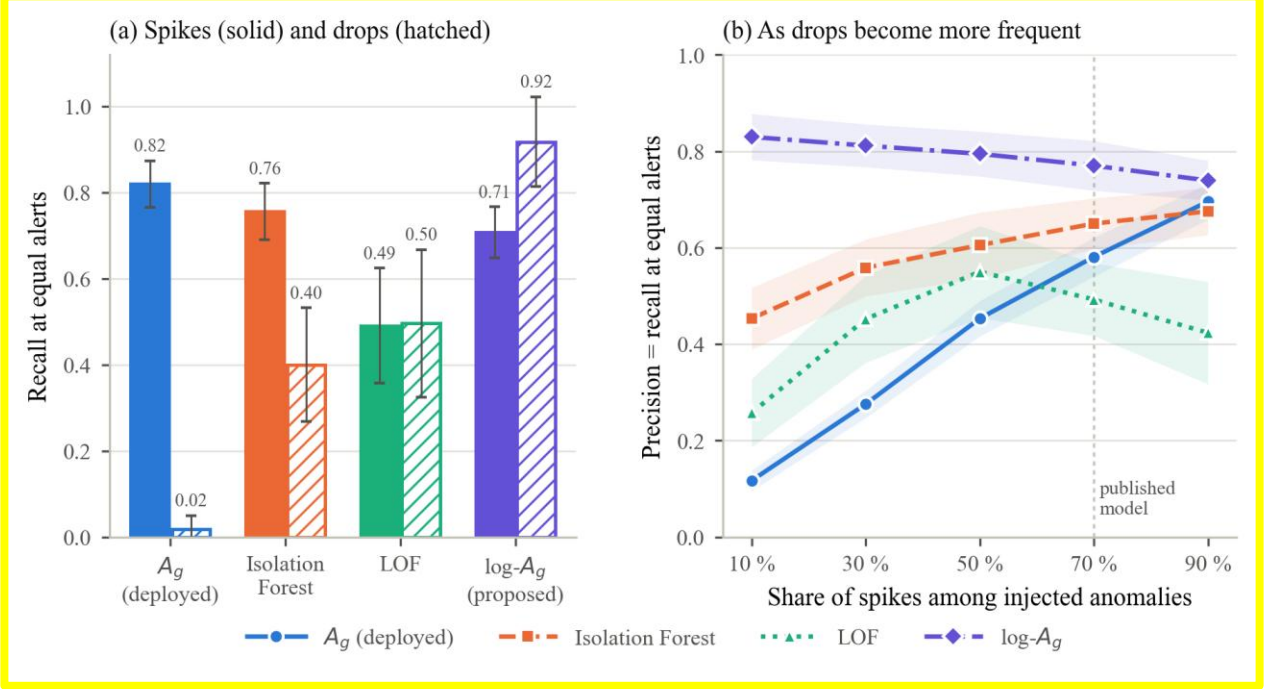


Fig. 2. Recall on injected spikes and drops under the published setting (a), and precision at the equalised operating point as the share of spikes varies (b). Mean over 20 seeds; bars and bands show one standard deviation.

On a logarithmic scale a multiplicative drop and a multiplicative spike lie at symmetric distances from the mean, and log-A_g removes the blindness: under the published setting it finds 71 % of spikes and 92 % of drops, raising precision from 0.581 to 0.770 and the area under the precision–recall curve from 0.677 to 0.860, above Isolation Forest (0.650 and 0.728). The advantage holds across prevalences and compositions (Table 6) and fades as the anomalies weaken: at moderate magnitude it only ties Isolation Forest (0.398), and at subtle magnitude every detector stays below 0.20, where the limit is the magnitude of the departure rather than the scale of the score. The comparison is not neutral, however: the synthetic consumption is log-normal by construction, which makes log-A_g the matched statistic, so its synthetic margin is an upper bound. Section 5.6 confronts both scores with the real distribution.

Table 6. Precision at the equalised operating point (equal to recall and F1) under each injection setting, mean over 20 seeds; standard deviations stay between 0.015 and 0.124. Published setting: 3 % prevalence, 70 % spikes, strong magnitude.

Setting	A_g	Isolation Forest	LOF	log-A_g
Published setting	0.581	0.650	0.493	0.770
Prevalence 1 %	0.491	0.544	0.600	0.709
Prevalence 10 %	0.749	0.763	0.303	0.834
10 % spikes, 90 % drops	0.117	0.453	0.257	0.830
50 % spikes, 50 % drops	0.453	0.605	0.550	0.795
90 % spikes, 10 % drops	0.696	0.676	0.423	0.739
Moderate magnitude	0.365	0.398	0.304	0.398
Subtle magnitude	0.181	0.183	0.122	0.173

5.6 The Deployed Ranking on the Real Data

The arithmetic bound does not depend on the synthetic generator, so it can be checked on the warehouse itself. Replicating the deployed computation over the 1,553 borough–neighbourhood totals of the published export, with no filter, shows totals

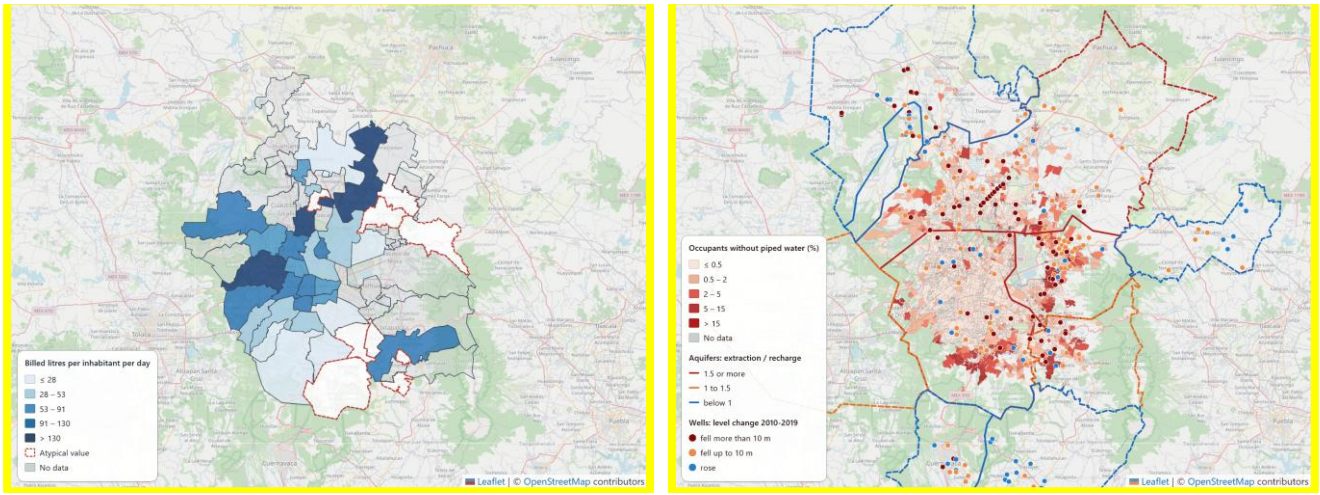
that are strongly right-skewed within every borough but one (median skewness 1.98) and much less skewed in logarithms (median -0.76). The ceiling μ/σ for a neighbourhood below its borough mean has a median of 0.85 across boroughs (0.27 to 1.64), and the highest score any of the 1,022 below-mean neighbourhoods actually reaches is 1.46, whereas the fifteenth entry the interface lists scores 3.85. Not one of the 15 neighbourhoods the module shows, nor of the top 3 % (47), lies below its borough mean: as deployed, the module can only ever report excess consumption.

The log-scale score does not settle the matter on real data; it inverts it. Because log totals are left-skewed, their long tail is the low one, and all 15 neighbourhoods at the top of the log-A_g ranking, and 45 of the top 47, lie below their borough mean; the two top-15 lists share no neighbourhood. A further 29 neighbourhoods have no billed volume at all over the semester, 27 of them with records in all three bimesters, concentrated in Tlalpan (12), Iztapalapa (6) and La Magdalena Contreras (5). Neither score treats them sensibly: A_g ranks them low and their logarithm is undefined. On these data no single symmetric score serves both directions. The design the evidence supports is three lists: excess and deficit, each ranked by log-A_q within its own direction, and neighbourhoods without billed volume as a separate flag, which in a source that bills unmetered intakes by averages (Section 3.4) may signal a data gap as much as an absence of use, which is why earlier spatial studies exclude colonias with zero consumption or no metering (Ramos-Bueno et al., 2021). The analysis is reproducible with `evaluacion-ml/ranking_real.py`.

5.7 From the City to the Metropolitan Area

The open export stops at the city boundary, but the water system does not. We therefore extended the warehouse additively to the official metropolitan delimitation of 2020 (CONAPO, 2025b): 63 municipalities, of which 16 are boroughs, 45 lie in the State of Mexico and two in Hidalgo. Sixty-two files, each checked against a manifest of sizes and hashes before loading, now feed 2.76 million rows in 43 tables, joined to the original location dimension through a conformed municipal dimension; the original fact tables are untouched. The new facts cover municipal services and billing (Instituto Nacional de Estadística y Geografía, n.d.), water revenue (SHCP, 2026), aquifers and wells (CONAGUA, 2026a), dams (CONAGUA, 2026b), drought (SMN, 2026), water quality (CONAGUA, n.d.), marginalisation (CONAPO, 2025a), land motion (Khorrami et al., 2023) and 555,752 citizen reports (SACMEX, n.d.).

Integration exposes what no single source shows. Municipal volumes cannot be compared without a filter: of the 47 municipalities outside Mexico City, 24 report no volume, 7 fall outside a plausibility band of 10 to 400 litres per inhabitant per day and 16 remain (Fig. 3a), while the revenue they declare does track federal records (Spearman 0.934 over 89 municipalities of the State of Mexico, Pearson 0.322), which confines the errors to a few captures. The reports also bound what can be asked of 2019: the source records no leaks before 2021, so all 39,685 reports of the semester are water shortages. Across 1,508 neighbourhoods, billed volume barely ranks shortage reports (Spearman 0.18), and the 435 neighbourhoods billed below 10 litres file almost as many reports per thousand inhabitants as the rest (4.15 against 4.93). Near-zero billing therefore does not mean near-zero demand, which supports the reading of Section 5.3. So does access: about 238,000 people in the metropolitan area live in dwellings without piped water (Fig. 3b), and in Mexico City the five boroughs with the largest share of them are among the six with the lowest billed volume. Groundwater gives the physical context. The four aquifers beneath the urban area extract between 1.49 and 1.96 times their recharge, in wells read in both years the mean water table fell between 6.0 and 18.7 metres from 2010 to 2019, and mean land motion reaches -23.5 cm per year in Nezahualcóyotl.



(a) Billed litres per inhabitant per day. *(b) Occupants without piped water, aquifers and wells.*

Fig. 3. Two views of the metropolitan map, each the output of SQL queries. Base map © OpenStreetMap contributors (ODbL).

The static deployment follows the warehouse without a server. Each layer of the map (https://gabrielhuav.github.io/Data_Warehouse_ZMVM/ijcopi/mapa-zm.html) is the literal result of one of thirteen SQL queries distributed in the repository, exported to `JSON` with its hash, and `verificar_mapa.py` repeats every query and compares the result byte by byte with the published files. A warehouse rebuilt from a clean clone returned the same bytes for all thirteen queries.

5.8 Performance Metrics

Lighthouse 13.4.1, run against the permanent static deployment, scores 89 on emulated desktop (first contentful paint 0.2 s, largest 0.3 s, total blocking time 200 ms) and 76 on emulated mobile over throttled 4G; accessibility, best practices and SEO score 100 on both. The remaining bound is structural: the static version downloads the full 3.9 MB export and aggregates on the client, so what a server computes once is computed by each visitor.

Query Performance over the Total Data Volume. Frontend metrics do not characterise the retrieval layer, so four analytical workloads were timed directly against the warehouse over all 70,886 fact rows: a paginated retrieval, a full-scan aggregation by borough, a full-scan aggregation by neighbourhood and bimester —the heaviest query the dashboard issues— and a top-*N* ranking. Each ran 200 times after ten warm-up calls on twelve logical cores with 32 GB of RAM and PostgreSQL 16.14 in a container, which is what makes the reported percentile meaningful. These are SQL latencies measured at the database, excluding serialisation and transport.

The profile, reported in Table 7, is the expected one for a star schema of this size: the three workloads that resolve through the dimensions stay near 25 ms while scanning the full fact table, whereas the aggregation by neighbourhood and bimester costs four times more because it groups over 1,553 locations rather than sixteen. The heaviest interactive query returns in about a tenth of a second, and the gap between mean and p95 stays below 25 % throughout, indicating that the cost is dominated by the scan rather than by contention.

Table 7. Response time of the retrieval workloads over the complete warehouse (70,886 fact rows), 200 executions after ten warm-up calls.

Workload	Mean (ms)	p95 (ms)	Executions
Paginated retrieval (`limit=100`)	32.3	35.2	200
Full aggregation by borough	24.1	27.0	200
Aggregation by neighbourhood and bimester	101.3	125.4	200
Top-10 neighbourhoods with ordering	23.1	23.6	200

5.9 Functional Tests

The operation of the ETL processes, the persistence in PostgreSQL, the parameterized queries and the information retrieval was verified. The tests included queries parameterized by year, bimester, borough, and neighbourhood; aggregated retrieval (total and average consumption by neighbourhood); generation of JSON responses; geospatial retrieval for visualization; and pagination with configurable limits. All modules produced results consistent with the expected structure. The dimensional model thus retrieves territorial information through combined filters of time, location and development index over the 70,886 records, supporting consumption queries by neighbourhood and borough, temporal aggregation, coordinate retrieval, geospatial exploration and ordering by total consumption.

6 Discussion

Against the closest antecedents the contribution is one of *provenance and preservation* rather than of scale. Soares et al. (2019) and Abecker et al. (2014) begin from data an organisation already controls; this work begins from a public export with no schema, versioning or update cadence, which moves the difficulty from modelling to consolidation and is why the cleaning criteria and their yield are reported explicitly. The second difference is durability: those systems depend on a running service, whereas here interface, graph and data are preserved as static artefacts any reader can re-execute (Hill et al., 2023; Sandve et al., 2013).

The practical implication is a caution about how the totals should be read. The boroughs that consume most are not those that consume most per inhabitant: Cuauhtémoc and Miguel Hidalgo concentrate commercial activity, and the warehouse supports the distinction because the source separates domestic from non-domestic billing —the non-domestic share reaches 32.5 % in Cuauhtémoc against 20.3 % in Gustavo A. Madero, twelve points apart between boroughs whose totals differ by barely a

third. Total consumption is therefore not a per-capita indicator and should not be used as one for allocating supply interventions; the per-account average, also stored, is the appropriate measure. The per-inhabitant figures of Section 5.3 confirm the reordering and reveal, in addition, how unevenly volumetric billing covers the city.

7 Conclusions and Future Work

7.1 Conclusions

A data warehouse was presented for the territorial information retrieval of water consumption in Mexico City, built from open sources through a reproducible ETL process and published with a continuous-deployment methodology that guarantees a permanent functional demonstration. The solution consolidated 70,886 consumption records and 1,553 locations across the 16 boroughs into a dimensional star schema on PostgreSQL.

The findings reveal that water consumption in Mexico City is concentrated in the boroughs of Cuauhtémoc and Miguel Hidalgo (Table 3) in billed totals, although Benito Juárez leads per inhabitant and volumetric billing captures little of the consumption of the eastern and southern periphery (Table 4), and that it varies across the available bimesters without establishing in this work a causal relationship with dry or rainy periods.

The three central contributions —the consolidation ETL, the validated graph layer derived from it, and the always-accessible demo methodology— constitute a replicable method for bringing heterogeneous official open data to queryable and verifiable territorial information. The extension adds a negative result with a practical consequence: as deployed, the atypical-consumption module can only report excess consumption; a log-scale score removes that blindness in synthetic tests, but on the real data it must be presented as separate lists for excess and deficit. Extending the warehouse to the 63 municipalities of the metropolitan area shows that municipal volumes need a quality filter before any comparison, and that low billing cannot be read as low demand.

7.2 Threats to Validity and Limitations

Six limitations bound what may be concluded. The temporal extent is the first: the source publishes three bimesters of a single year, so the seasonal variation observed is a within-semester comparison and the climate correlation rests on three aggregated points, enough to describe a direction but not to estimate one. The second concerns the anomaly module, whose evaluation uses synthetic perturbations aligned with the detectors; its figures are an upper bound and say nothing about real leaks. The sensitivity analysis widens the injection model but not its family: anomalies remain multiplicative over a log-normal baseline, which favours the log-scale score. The third is geometric: neighbourhoods are represented by centroids rather than polygons, so the adjacency projected as `agua:nearbyWithin1500m` is proximity within 1.5 km and not shared boundaries. The fourth is descriptive: the association between development index and consumption is an observation over the warehouse, and attributing it to a socioeconomic mechanism would require an external source the warehouse does not integrate. The fifth is the billing basis and its coverage: values mix metered and averaged billing in a proportion the export does not disclose, and although they appear to cover nearly all volumetric billing, they add up to about 38 % of published per-capita consumption, with the gap concentrated in the periphery. Territorial contrasts and atypicality scores therefore describe billing, and coincide with physical use only where intakes are metered and billed by volume. The sixth concerns the metropolitan layers: they combine reference years from 2018 to 2020, the neighbourhood polygons are CONAPO's rather than the utility's, several joins rest on spatial proximity, and reports count what citizens file, not what is missing.

7.3 Future Work

The most immediate line is to materialise the knowledge graph into a persistent triple store with a public SPARQL endpoint, and to replace the centroid-based adjacency with official neighbourhood polygons, which would make the topological predicates exact. Aligning the territorial identifiers with external authorities would let the resource be joined with other municipal datasets without prior schema agreement, which is the interoperability argument that motivates the graph layer. The anomaly module should be re-evaluated when labelled leak data exist, with the three-list design of Section 5.6 implemented in the interface. A data cube over the extended warehouse is the next step for the metropolitan analysis. Finally, the method transfers to other open urban-data domains —mobility, air quality and seismic risk— where the same combination of dispersed official files and absent query tooling recurs.

References

Abecker, A., Brauer, T., Magoutas, B., Mentzas, G., Papageorgiou, N., & Quenzer, M. (2014). A sensor and semantic data warehouse for integrated water resource management. In *EnviroInfo 2014*. BIS-Verlag.

Berro, A., Megdiche, I., & Teste, O. (2015). A content-driven ETL processes for open data. In N. Bassiliades et al. (Eds.), *New trends in database and information systems II*. Springer. https://doi.org/10.1007/978-3-319-10518-5_3

Breunig, M. M., Kriegel, H.-P., Ng, R. T., & Sander, J. (2000). LOF: Identifying density-based local outliers. In *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data* (pp. 93–104). ACM. https://doi.org/10.1145/342009.335388

Car, N. J., & Homburg, T. (2022). GeoSPARQL 1.1: Motivations, details and applications of the decadal update to the most important geospatial LOD standard. *ISPRS International Journal of Geo-Information, 11*(2), 117. https://doi.org/10.3390/ijgi11020117

CONAGUA [Comisión Nacional del Agua]. (2026a). *Aguas subterráneas* [Data set]. datos.gob.mx. https://www.datos.gob.mx/dataset/aguas_subterraneas

CONAGUA [Comisión Nacional del Agua]. (2026b). *Sistema de Información Hidrológica: Series diarias de presas* [Data set, 15 CSV files]. Retrieved October 30, 2026, from https://sih.conagua.gob.mx/basedatos/Presas/VBRMX.csv

CONAGUA [Comisión Nacional del Agua]. (n.d.). *Indicadores de calidad del agua* [Data set]. Retrieved October 31, 2026, from https://www.gob.mx/conagua/es/articulos/indicadores-de-calidad-del-agua?idiom=es

CONAPO [Consejo Nacional de Población]. (2025a). *Índices de marginación* [Data set]. datos.gob.mx. https://www.datos.gob.mx/dataset/d8f2a534-bcee-4114-853d-82982a81ce24

CONAPO [Consejo Nacional de Población]. (2025b). *Las metrópolis de México, 2020* [Data set]. datos.gob.mx. https://www.datos.gob.mx/dataset/da65ccd9-888c-4ce8-b3b5-6bd867be344d

Hill, D., Dunham, C., Larsen, D. A., & Collins, M. (2023). Operationalizing an open-source dashboard for wastewater-based surveillance. *MethodsX, 11*, 102299. https://doi.org/10.1016/j.mex.2023.102299

INFO CDMX [Instituto de Transparencia, Acceso a la Información Pública, Protección de Datos Personales y Rendición de Cuentas de la Ciudad de México]. (2022). *Resolución del recurso de revisión INFOCDMX/RR.IP.0679/2022* (Sujeto obligado: Sistema de Aguas de la Ciudad de México). https://documentosobligaciones.infocdmx.org.mx/ST/LTAIPRC-2016-OT/Art133/Fr01/2022/T02-INFOCDMX.RR.IP.0679.2022.pdf

Instituto Nacional de Estadística y Geografía. (2021). *Censo de Población y Vivienda 2020: Principales resultados por localidad (ITER), Ciudad de México* [Data set, CSV]. INEGI. https://www.inegi.org.mx/programas/ccpv/2020/

Instituto Nacional de Estadística y Geografía. (n.d.). *Censo Nacional de Gobiernos Municipales y Demarcaciones Territoriales de la Ciudad de México, 2019 and 2021 editions* [Data set]. Retrieved October 31, 2026, from https://www.inegi.org.mx/programas/cngmd/2021/

Khorrami, M., Shirzaei, M., Ghobadi-Far, K., Werth, S., Carlson, G., & Zhai, G. (2023). Groundwater volume loss in Mexico City constrained by InSAR and GRACE observations and mechanical models. *Geophysical Research Letters, 50*(5), e2022GL101962. https://doi.org/10.1029/2022GL101962

Kimball, R., & Ross, M. (2013). *The data warehouse toolkit: The definitive guide to dimensional modeling* (3rd ed.). Wiley.

Lippolis, A. S., Lodi, G., & Nuzzolese, A. G. (2025). The Water Health Open Knowledge Graph. *Scientific Data, 12*, 274. https://doi.org/10.1038/s41597-025-04537-4

Liu, F. T., Ting, K. M., & Zhou, Z.-H. (2008). Isolation forest. In *2008 Eighth IEEE International Conference on Data Mining* (pp. 413–422). IEEE. https://doi.org/10.1109/ICDM.2008.17

Medina-Rivas, C. M., Morales-Novelo, J. A., Rodríguez-Tapia, L., & Revollo-Fernández, D. A. (2025). Mexico city's decline in per capita domestic water use: A comprehensive spatial-temporal study. *Urban Water Journal, 22*(1), 1–15. https://doi.org/10.1080/1573062X.2024.2423400

Medina-Rivas, C. M., Rodríguez-Tapia, L., Morales-Novelo, J. A., & Revollo-Fernández, D. A. (2022). Spatial inequality of domestic water consumption in Mexico city. *Water Resources and Economics, 40*, 100210. https://doi.org/10.1016/j.wre.2022.100210

Purves, R. S., Clough, P., Jones, C. B., Hall, M. H., & Murdock, V. (2018). Geographic information retrieval. *Foundations and Trends in Information Retrieval, 12*(2–3), 164–318. https://doi.org/10.1561/1500000034

Ramos-Bueno, A., Galeana-Pizaña, J. M., & Perevochtchikova, M. (2024). Urban water consumption analysis through a spatial panel modeling approach: A case study of Mexico City, 2004–2022. *Water Supply, 24*(9), 3179–3195. https://doi.org/10.2166/ws.2024.191

Ramos-Bueno, A., Perevochtchikova, M., & Chang, H. (2021). Socio-spatial analysis of residential water demand in Mexico City. *Tecnología y Ciencias del Agua, 12*(2), 59–110. https://doi.org/10.24850/j-tyca-2021-02-02

Roussey, C., Guéret, C., & Laporte, M.-A. (2023). Editorial: Knowledge graph technologies—The next frontier of the food, agriculture, and water domains. *Frontiers in Artificial Intelligence, 6*, 1319844. https://doi.org/10.3389/frai.2023.1319844

SACMEX. (2026, August 28). *Consumo de agua: Consumo Agua (1° Semestre 2019)* [Data set, CSV, with data dictionary and clarifying note]. Portal de Datos Abiertos de la Ciudad de México, Agencia Digital de Innovación Pública. Retrieved October 18, 2026, from https://datos.cdmx.gob.mx/dataset/consumo-agua

SACMEX. (n.d.). *Reportes de agua* [Data set, CSV]. Portal de Datos Abiertos de la Ciudad de México. https://datos.cdmx.gob.mx/dataset/reportes-de-agua

Saito, T., & Rehmsmeier, M. (2015). The precision-recall plot is more informative than the ROC plot when evaluating binary classifiers on imbalanced datasets. *PLoS ONE, 10*(3), e0118432. https://doi.org/10.1371/journal.pone.0118432

Sandve, G. K., Nekrutenko, A., Taylor, J., & Hovig, E. (2013). Ten simple rules for reproducible computational research. *PLoS Computational Biology, 9*(10), e1003285. https://doi.org/10.1371/journal.pcbi.1003285

SHCP [Secretaría de Hacienda y Crédito Público]. (2026). *Derechos de agua municipal* [Data set]. datos.gob.mx. https://www.datos.gob.mx/dataset/derechos_agua_municipal

SMN [Servicio Meteorológico Nacional]. (2026). *Monitor de Sequía de México: Municipios* [Data set]. CONAGUA. https://smn.conagua.gob.mx/tools/RESOURCES/Monitor%20de%20Sequia%20en%20Mexico/MunicipiosSequia.xlsx

Soares, J., Leite, P., Teixeira, P., Lopes, N., & Silva, J. P. (2019). Data warehouse for the monitoring and analysis of water supply and consumption. In I. Ramos, R. Quaresma, P. Resende da Silva, & T. Oliveira (Eds.), *Information systems for Industry 4.0*. Springer. https://doi.org/10.1007/978-3-030-14850-8_1

Vassiliadis, P. (2009). A survey of extract–transform–load technology. *International Journal of Data Warehousing and Mining, 5*(3), 1–27. https://doi.org/10.4018/jdwm.2009070101

Velázquez Arrieta, E. U., Pulido Morales, O. F., García López, E., Hernández Martínez, C. A., & Hurtado Avilés, G. (in press). Territorial information retrieval from heterogeneous open data through the construction of a data warehouse for water management in Mexico City. In *Proceedings of the 5th International Conference on Ontologies and Knowledge Graphs (ICOKG 2026)*. Springer.

Vetrò, A., Canova, L., Torchiano, M., Orozco Minotas, C., Iemma, R., & Morando, F. (2016). Open data quality measurement framework. *Government Information Quarterly, 33*(2), 325–337. https://doi.org/10.1016/j.giq.2016.02.001

Villa Vargas, J. M., Hurtado Avilés, G., & Climent Hernández, J. A. (2026). Cuando México tiembla: La historia contada por los datos. *AZCATL Revista de Divulgación en Ciencias, Ingeniería e Innovación, 4*(6), 28–33. https://doi.org/10.24275/AZC2026E1004

World Wide Web Consortium. (2012, September 27). *R2RML: RDB to RDF mapping language* [W3C Recommendation]. https://www.w3.org/TR/r2rml/

World Wide Web Consortium. (2014, January 16). *The RDF Data Cube vocabulary* [W3C Recommendation]. https://www.w3.org/TR/vocab-data-cube/

Data Availability Statement: The source files, the warehouse, the R2RML mapping, the scripts and the static demonstration are distributed in the archived release https://github.com/gabrielhuav/Data_Warehouse_static/releases/tag/v1.5-icokg2026. The extended warehouse, its verified manifest of sources, the SQL queries and the metropolitan map are in the repository https://github.com/gabrielhuav/Data_Warehouse_ZMVM.

Declaration on the Use of Generative AI: The authors used a generative AI assistant, under their direction, for language editing, LaTeX formatting, converting the manuscript to the journal template, searching and verifying the regional data sources, and writing the ETL, queries and map of the metropolitan extension, drafting text subsequently revised and approved by the authors, and writing the auxiliary scripts distributed with the artefact. The research questions, the dimensional design, the original ETL, the data collection of the conference version, the experiments and the interpretation of the results are the authors' own work. Every reported quantity was recomputed from the published source files and verified by the authors, and the scripts that perform that recomputation are distributed with the release. The authors take full responsibility for the content.

Conflicts of Interest: The authors declare that they have no conflicts of interest relevant to the content of this article.